



OUTLIERS

Judgement

Layer 5 of 8 · What your system cannot
work out on its own.

Something has to weigh it up


Your system can now tell you what happened. It cannot tell you what any of it means. That needs judgement, and this is the first layer that has any.

WHAT THE LAST LAYER LEFT YOU WITH

You have a history that fills itself. What you do not have is anything that reads it and forms a view. A program can tell you that somebody replied. It cannot tell you whether the reply was warm, whether the person is worth your time, or what to say next.

WHAT THIS LAYER BUILDS

Assistants. Each one is given a single job, its own written instructions, and permission to use particular things and nothing else. They read what your system knows, form a view, and hand it to you.



How it works

An assistant is not magic and should not be treated as such. It is a worker you briefed, and you judge it the same way.

WHAT AN ASSISTANT ACTUALLY IS

A set of written instructions handed to an AI, the same kind you would type a question into, plus one clearly defined job, plus a short list of actions it is allowed to take and nothing outside that list. That is the whole of it. The instructions are a text file you can read and change, so you can see exactly what you asked for.

JUDGE ONE THE WAY YOU WOULD JUDGE A PERSON

You already know what good and bad look like when a human does a job. If you hired someone to go through your contacts and tell you which are worth a call, you would know within ten minutes whether they had understood. Use exactly that instinct here. If the output would annoy you coming from a person, it is not acceptable coming from an assistant either.

ONE JOB EACH, NEVER ONE BIG ONE

The tempting thing is a single assistant that does everything. It cannot be tested in parts, and correcting one behaviour changes another you never touched. Several small ones can each be checked, improved or thrown away on their own.

EVERY FACT SAYS WHERE IT CAME FROM

An assistant that is right and an assistant that has invented something produce output that looks identical. Same confidence, same tidy sentences. A week later you cannot tell which is which by reading it.

So nothing gets written into your records without saying where it came from and when. If the evidence is thin, it is not saved at all: it waits in a list for you to look at. Your records stay a record of things that are true, and everything uncertain stays visibly uncertain.

Build it

One command. It asks you a few questions so what you get is yours, then does the work.

```
git clone https://github.com/OUTLIERS-ai/outliers-crm-05-judgement
cd outliers-crm-05-judgement
python install.py
```

This layer needs Layers 1 to 4. There has to be something in your history for an assistant to read.

WHAT IT ASKS YOU, AND WHY

QUESTION	WHAT IT CHANGES
What makes someone a good fit for you, in your own words?	Written as you would say it out loud. The assistant is given your definition, not a generic one, so it is scoring for your business.
What would make you drop someone?	Just as useful as the positive version, and usually easier to answer honestly.

WHAT YOU END UP WITH

- A worked assistant with one job, whose instructions you can read and edit.
- The rule that everything written carries its source, applied automatically.
- A holding list where anything uncertain waits for you instead of entering your records.

Now use it

Ten minutes, and it will tell you something about your own contacts.

Let the assistant score the people in your CRM against the definition you gave it.

Then find one you disagree with, and read the reason it gave. One of two things is true: either it misread something, which tells you the instructions need a sentence adding, or it spotted something you had not, which is the whole reason you built it.

IF YOU ALREADY HAD SOMETHING LIKE THIS

Your test: take one thing your current assistant or tool has told you and try to find out where the information came from. If you cannot trace it back to a source, you do not know whether it is true.

WHAT THIS LAYER LEAVES UNSOLVED

You now have several assistants working, and any one of them could put a message in front of a real human being. Nothing is stopping two of them acting at once, or one of them checking its own work and passing it. That needs control which does not depend on them behaving well.

NEXT: LAYER 6 • ORDER AND SAFETY

Sequence, limits, and the rule that nothing reaches another person without your hand on it. Built into the machinery rather than written in the instructions.

WORDS USED IN THIS GUIDE

AI model	The kind of AI you type a question into and get an answer back from.
Assistant	Written instructions, plus one clearly defined job, plus a short list of actions it may take.
Tool	One specific action an assistant is allowed to take. Not an app.
Source	Where a fact came from, recorded next to the fact itself.
Holding list	Where uncertain things wait for you, instead of being saved as true.
Staging	Where anything uncertain waits for you, instead of being saved as true.